

```
In [83]: #make sure to download all of these packages
#python version that I used was 3.10, but most version of 3 should work
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
import seaborn as sns
%matplotlib inline

from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
```

```
In [84]: #gets the values that are outside of the IQR for a specific set features (classes =
def outside_of_iqr(df : pd.DataFrame, investigate: str, classes: str) -> dict:

    #get all unique values of the class
    values = df[classes].unique()
    dictionary = dict()

    for i in values:
        specific_quality = df.loc[df[classes] == i]

        #get the IQR of the feature that is a certain class
        q1 = specific_quality[investigate].quantile(0.25)
        q3 = specific_quality[investigate].quantile(0.75)

        #get temporary values of total ones that are outside of the range of a part
        temp = specific_quality[investigate].loc[(specific_quality[investigate] > q3) | (specific_quality[investigate] < q1)]
        dictionary.update({i: len(temp)})

    return dictionary
```

```
In [85]: red_wine_data = pd.read_csv("data/winequality-red.csv",delimiter=";")
white_wine_data = pd.read_csv("data/winequality-white.csv",delimiter=";")
```

```
In [86]: red_wine_data
```

Out[86]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.51
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.61
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.61
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.51
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.51
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.51
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.71
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.71
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.61

1599 rows × 12 columns

In [87]: red_wine_data.describe()

Out[87]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.000000
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.000000
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

Seems to be no missing data within all of the features at least for red

```
In [88]: fig = go.Figure()
#getting the individual amounts of times the quality(labels of the wine) appears in the data
quality_amounts_red = red_wine_data["quality"].groupby(red_wine_data["quality"]).count()
quality_amounts_white = white_wine_data["quality"].groupby(white_wine_data["quality"]).count()
```

```

#adding to a graph
fig.add_trace(go.Bar(x=quality_amounts_red.index,y=quality_amounts_red.values,name=
fig.add_trace(go.Bar(x=quality_amounts_white.index,y=quality_amounts_white.values,n

fig.update_layout(title="Quantity of Quality")
fig.show()

```

Overall the data set is imbalanced, so a accuracy measure will have to handle imbalanced.
There is much more data in white, but at least it appears that they follow similar distribution

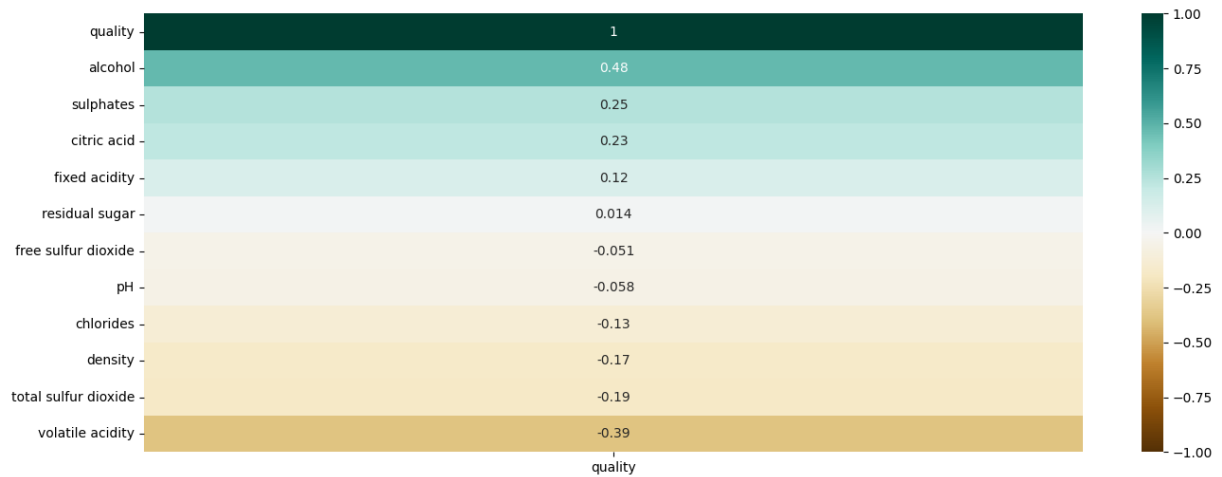
```

In [89]: """
vmin,vmax - the range of values for colormap(min-max)
cmap - sets the specific colormap to use
cetner - takes a float to centera color map
annot - if True sets the correlation values to appear
cbar - if False, the colorbar disapears

"""
plt.figure(figsize=(16,6))
sns.heatmap(red_wine_data.corr()[["quality"]].sort_values(by='quality',ascending=Fa

```

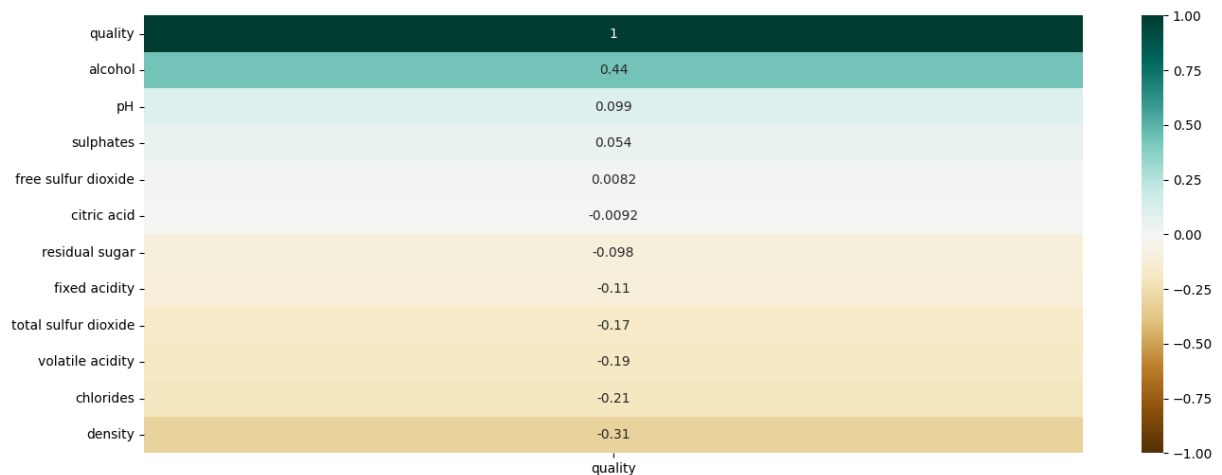
Out[89]: <AxesSubplot: >



So, this diagram shows the relationship (correlation) between a feature and our label (quality). If the value is high(1) or low(-1) that entails that there is a strong correlation between the label and the quality. Unfortunately we don't have many high ones of the bat, high ones being (volatile acidity and alcohol). However, this is raw data that we can try read the data a bit more (binning and such).

```
In [90]: plt.figure(figsize=(16,6))
sns.heatmap(white_wine_data.corr()[["quality"]].sort_values(by='quality',ascending=
```

```
Out[90]: <AxesSubplot: >
```



In white wines it appears that the lowest correlations (between -0.1 and 0.1) are the same (free sulfur dioxide, residual sugar, pH), so it seems fair to drop these from consideration; however it might be worthwhile to check later if there are some values that are throwing off others. For initial EDA these will be dropped.

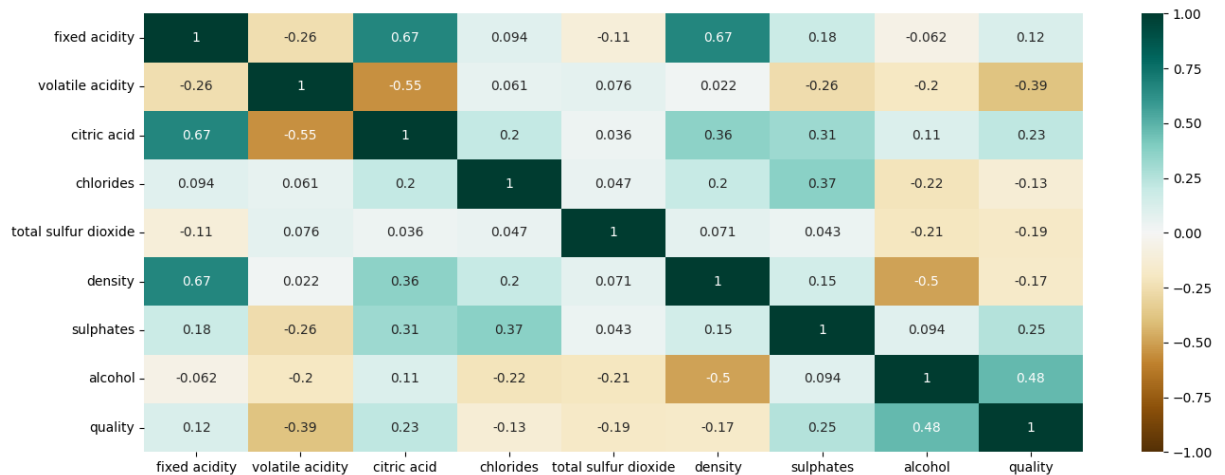
Citric acid and sulphates appeared to show strong correlation in red but not white. Will need to combine data sets and do EDA on all at the same time

Overall, it seems that red wine has much stronger correlation between specific attributes; therefore we are going to use the red wine for our models. This comes with an unfortunate

tradeoff, as the white wine dataset has more data overall, which might pose issues with our neural network attempt.

```
In [91]: red_wine_data = red_wine_data.drop(["free sulfur dioxide","residual sugar" ,"pH"],a
```

```
In [92]: """
Overall shows the direct or indirect correlation between features of a dataset.
of two features. Measuring anything beyond the relationship of two is not possi
"""
plt.figure(figsize=(16, 6))
heatmap = sns.heatmap(red_wine_data.corr(), vmin=-1, vmax=1, annot=True, cmap='BrBG
```



Attribute: Alcohol

```
In [93]: #Alcohol stuff
from plotly.subplots import make_subplots

fig = make_subplots(rows=2,cols=2)

#histogram for red
fig.append_trace(go.Histogram(
    x=red_wine_data["alcohol"],
    name="red hist"), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=red_wine_data["alcohol"],
    name="red box"
),row=1,col=2)

#histogram for white
fig.update_layout(height=600, width=1200, title_text="Alocohol data")
fig.show()
```



Appears to be a bit skewed but overall, not a huge amount of problems.

```
In [94]: px.box(red_wine_data,x="quality",y="alcohol",title="Alchol vs Quality")
```

Can see a bit of a trend here, as the value of alcohol increases so does the actual value of alcohol, and it appears that there is a considerable jump from 5 to 6

```
In [95]: mean_red = red_wine_data.loc[(red_wine_data["quality"] >= 7)][["alcohol"]].mean()
greater_than_mean = red_wine_data.loc[red_wine_data["quality"] < 7][["alcohol"]].valu

less_than_mean = red_wine_data.loc[(red_wine_data["quality"] >= 7)][["alcohol"]].valu
sum(less_than_mean), sum(greater_than_mean), len(red_wine_data)
```

Out[95]: (107, 140, 1599)

It seems that we might want to split the classes into just two binary variables, overall it seems that our bet will to have (1-5) and (6-10)

```
In [96]: red_wine_data["alcohol_higher"] = 0
red_wine_data.loc[red_wine_data["alcohol"] >= mean_red, "alcohol_higher"] = 1
```

Attribute: density

```
In [97]: #density seems to be very very small differences, I doubt this will be any help but  
red_wine_data["density"].describe()
```

```
Out[97]: count      1599.000000  
mean         0.996747  
std          0.001887  
min          0.990070  
25%          0.995600  
50%          0.996750  
75%          0.997835  
max          1.003690  
Name: density, dtype: float64
```

```
In [98]: px.histogram(red_wine_data,x="density")
```

```
In [99]: px.box(red_wine_data,x="quality",y="density",title="Red Wine")
```


almost a perfect distribution for normal

Attribute: volatile acidity

```
In [100]: red_wine_data["volatile acidity"].describe()
```

```
Out[100]: count    1599.000000
          mean      0.527821
          std       0.179060
          min       0.120000
          25%       0.390000
          50%       0.520000
          75%       0.640000
          max       1.580000
          Name: volatile acidity, dtype: float64
```

There is pretty good standard deviation, so plausible that we might be able to use this data pretty well for banding.

```
In [101]: fig = make_subplots(rows=1, cols=2)
```

```
#histogram for red
fig.append_trace(go.Histogram(
    x=red_wine_data["volatile acidity"],
    name="red hist"), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=red_wine_data["volatile acidity"],
    name="red box"
),row=1,col=2)

fig.update_layout(height=600, width=1200, title_text="Volatile Acidity")
fig.show()
```

Bit skewed but more normal if anything

In [102... `px.box(red_wine_data,x="quality",y="volatile acidity",title="Red Wine")`

```
In [103... means_of_red = red_wine_data.groupby(red_wine_data["quality"]).mean()  
medians_of_red = red_wine_data.groupby(red_wine_data["quality"]).median()
```

```
In [104... means_of_red["volatile acidity"],medians_of_red["volatile acidity"]
```

```
Out[104]: (quality  
3      0.884500  
4      0.693962  
5      0.577041  
6      0.497484  
7      0.403920  
8      0.423333  
Name: volatile acidity, dtype: float64,  
quality  
3      0.845  
4      0.670  
5      0.580  
6      0.490  
7      0.370  
8      0.370  
Name: volatile acidity, dtype: float64)
```

```
In [105... #getting tuples that have a quality that is greater than five  
greater_than_five = red_wine_data.loc[red_wine_data["quality"] > 5]
```

```
greater_than_five["volatile acidity"].describe()
```

```
Out[105]: count      855.000000
          mean        0.474146
          std         0.161999
          min         0.120000
          25%         0.350000
          50%         0.460000
          75%         0.580000
          max         1.040000
          Name: volatile acidity, dtype: float64
```

```
In [106... #getting the avlues that are greater than five, and have less than 0.6 volatile aci
high_quality_less = greater_than_five.loc[greater_than_five["volatile acidity"] < 0.6]
#getting the values that are greater than five, and have more than 0.6 volatile aci
high_quality_more = greater_than_five.loc[greater_than_five["volatile acidity"] >= 0.6]
print("Amount that fit the band " + str(len(high_quality_less)))
print("Amount that do not fit the band " + str(len(high_quality_more)))
```

Amount that fit the band 664

Amount that do not fit the band 191

```
In [107... less_than_five = red_wine_data[red_wine_data["quality"] < 5]
less_than_five["volatile acidity"].describe()
```

```
Out[107]: count      63.000000
          mean        0.724206
          std         0.247970
          min         0.230000
          25%         0.565000
          50%         0.680000
          75%         0.882500
          max         1.580000
          Name: volatile acidity, dtype: float64
```

```
In [108... high_quality_less = less_than_five.loc[less_than_five["volatile acidity"] <= 0.6]
high_quality_more = less_than_five.loc[less_than_five["volatile acidity"] > 0.6]
print("Amount that fit the band " + str(len(high_quality_more)))
print("Amount that do not fit the band " + str(len(high_quality_less)))
```

Amount that fit the band 42

Amount that do not fit the band 21

```
In [109... #creating a binary feature
red_wine_data["va_high"] = 0
red_wine_data.loc[red_wine_data["volatile acidity"] >= 0.6, "va_high"] = 1
red_wine_data
```

Out[109]:

	fixed acidity	volatile acidity	citric acid	chlorides	total sulfur dioxide	density	sulphates	alcohol	quality	alco
0	7.4	0.700	0.00	0.076	34.0	0.99780	0.56	9.4	5	
1	7.8	0.880	0.00	0.098	67.0	0.99680	0.68	9.8	5	
2	7.8	0.760	0.04	0.092	54.0	0.99700	0.65	9.8	5	
3	11.2	0.280	0.56	0.075	60.0	0.99800	0.58	9.8	6	
4	7.4	0.700	0.00	0.076	34.0	0.99780	0.56	9.4	5	
...
1594	6.2	0.600	0.08	0.090	44.0	0.99490	0.58	10.5	5	
1595	5.9	0.550	0.10	0.062	51.0	0.99512	0.76	11.2	6	
1596	6.3	0.510	0.13	0.076	40.0	0.99574	0.75	11.0	6	
1597	5.9	0.645	0.12	0.075	44.0	0.99547	0.71	10.2	5	
1598	6.0	0.310	0.47	0.067	42.0	0.99549	0.66	11.0	6	

1599 rows × 11 columns



This also might cause overfitting due to the correlation with volatile acidity in general. However, it should at least be helpful in determining red wines better. Check confusion matrix at the end and run models with and without the values

Attribute : Total sulfur dioxide

In [110... `red_wine_data["total sulfur dioxide"].describe(), white_wine_data["total sulfur dio`

```
Out[110]: (count      1599.000000
          mean       46.467792
          std        32.895324
          min         6.000000
          25%        22.000000
          50%        38.000000
          75%        62.000000
          max       289.000000
          Name: total sulfur dioxide, dtype: float64,
          count      4898.000000
          mean       138.360657
          std        42.498065
          min         9.000000
          25%       108.000000
          50%       134.000000
          75%       167.000000
          max       440.000000
          Name: total sulfur dioxide, dtype: float64)
```

```
In [111... px.histogram(red_wine_data,x="total sulfur dioxide")
```

```
In [112... px.box(red_wine_data,x="quality",y="total sulfur dioxide",title="Red Wine TSD")
```

This data is highly skewed and it might be worthwhile to try and to a transformation to smooth it out. Either smooth or turn into a standard scaler.

Attribute: citiric acid

```
In [113... red_wine_data["citric acid"].describe()
```

```
Out[113]: count    1599.000000
mean         0.270976
std          0.194801
min          0.000000
25%          0.090000
50%          0.260000
75%          0.420000
max          1.000000
Name: citric acid, dtype: float64
```

```
In [114... px.histogram(red_wine_data,x="citric acid")
```

In [115... `px.box(red_wine_data,x="quality",y="citric acid",title="Red Wine Citric acid")`


```
In [116... outside_of_iqr(red_wine_data, "citric acid", "quality")
```

```
Out[116]: {5: 327, 6: 309, 7: 96, 4: 25, 8: 9, 3: 6}
```

Even though it might look like there could be not real outliers, there is still a great deal of variance within the upper and lower, fence out outside of the IQR

Main winners: ones in parenthesis are simplified attributes. Might be helpful in certain algorithms, we can try to use both main attribute and other attribute, but should be wary of overfitting

Positive - Alcohol (alcohol_higher), sulphates, citric acid

Negative - Volatile acidity (va_high), Total sulfur dioxide, density

Possible drops -- fixed acidity, and chlorides

```
In [117... red_wine_data["graphing qualities"] = ""  
red_wine_data.loc[red_wine_data["quality"] > 6, "graphing qualities"] = "7-8"
```

```
red_wine_data.loc[red_wine_data["quality"] < 5, "graphing qualities"] = "3-4"
red_wine_data.loc[(red_wine_data["quality"] < 7) & (red_wine_data["quality"] > 4),
```

Chlorides/Sulphates : sulphates have a good correlation, but chlorides do not. However, both have a decent correlation together. It might be worth the time to remove chlorides to reduce confusion in the model, or integrate it somehow but lose unimportant information.

In [118...

```
#Alcohol stuff
fig = make_subplots(rows=2,cols=3)

#histogram for red
fig.append_trace(go.Histogram(
    x=red_wine_data["chlorides"],
    nbinsx=10,
    name="chloride hist"
), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=red_wine_data["chlorides"],
    name="chloride box"
),row=1,col=2)

fig.append_trace(go.Box(
    x=red_wine_data["quality"],
    y=red_wine_data["chlorides"],
    name="quality chlorides"
), row=1,col=3)

fig.append_trace(go.Histogram(
    x=red_wine_data["sulphates"],
    nbinsx=10,
    name="sulphate hist"
), row=2,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=red_wine_data["sulphates"],
    name="sulphate box"
),row=2,col=2)

fig.append_trace(go.Box(
    x=red_wine_data["quality"],
    y=red_wine_data["sulphates"],
    name="quality sulphates"
), row=2,col=3)

fig.update_layout(height=600, width=1250, title_text="Sulphates/Chlorides")
fig.show()
```



```
In [119... fig = px.scatter(red_wine_data, x="chlorides", y="sulphates", color="graphing quali  
fig.show()
```

Both have very similar distributions, and the correlation just appears to arise that they cluster around a similar area. There is no way to actually discern where certain values would be in their combination, so I would wager we could just drop chlorides.

fixed acidity/density/citric acid

In [120...

```
#Alcohol stuff
fig = make_subplots(rows=3,cols=3)

fig.append_trace(go.Histogram(
    x=red_wine_data["fixed acidity"],
    name="fixed acidity hist"
), row=1,col=1)

fig.append_trace(go.Box(
    x=red_wine_data["fixed acidity"],
    name="fixed acidity box"
),row=1,col=2)

fig.append_trace(go.Box(
    x=red_wine_data["quality"],
```

```
y=red_wine_data["fixed acidity"],
name="quality fixed acidity"
), row=1,col=3)

fig.append_trace(go.Histogram(
    x=red_wine_data["density"],
    name="density hist"
), row=2,col=1)

fig.append_trace(go.Box(
    x=red_wine_data["density"],
    name="density box"
),row=2,col=2)


fig.append_trace(go.Box(
    x=red_wine_data["quality"],
    y=red_wine_data["density"],
    name="quality density"
), row=2,col=3)

fig.append_trace(go.Histogram(
    x=red_wine_data["citric acid"],
    name="citric acid hist"
), row=3,col=1)

fig.append_trace(go.Box(
    x=red_wine_data["citric acid"],
    name="citric acid box"
),row=3,col=2)

fig.append_trace(go.Box(
    x=red_wine_data["quality"],
    y=red_wine_data["citric acid"],
    name="quality citric acid"
), row=3,col=3)

fig.update_layout(height=600, width=1250, title_text="Fixed Acidity/Density/Citric
fig.show()
```



Fixed acidity, and density have a similar distribution if you do not consider the scaling of the values. On the other hand, citric acid seems to be much different overall.

```
In [121... red_wine_data = red_wine_data.drop(["graphing qualities"],axis=1)
red_train, red_test = train_test_split(red_wine_data,test_size=0.2,stratify=red_wine_data["graphing qualities"])
```

When it comes to the analysis of individual features, there does not seem anything more we can do except scale and remove outliers when needed. For KNN we should scale, DT won't need much, and NN will need scaling of some kind. However, there is a chance that there are interactions between features in more complex way. If we had more time we could try and figure out, but I have no experience beyond direct correlation. Therefore, we can leave that to another day.

```
In [122... red_train.to_csv("data/red_wine_train.csv")  
red_test.to_csv("data/red_wine_test.csv")
```

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
import keras
import sklearn
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [2]: train = pd.read_csv('data/red_wine_train.csv', index_col=0)

#dropping columns based on heatmap correlations in EDA
train = train.drop(['fixed acidity', 'chlorides', 'density', 'total sulfur dioxide'])
train.head()
```

```
Out[2]:
```

	volatile acidity	citric acid	sulphates	alcohol	quality
1348	0.655	0.03	0.39	9.5	5
117	0.560	0.12	0.50	9.4	6
1150	0.330	0.32	0.76	12.8	7
235	0.630	0.00	0.58	9.0	6
91	0.490	0.28	1.95	9.9	6

```
In [3]: test = pd.read_csv('data/red_wine_test.csv', index_col=0)
test = test.drop(['fixed acidity', 'chlorides', 'density', 'total sulfur dioxide'],
```

```
In [4]: ###begin raw
```

```
In [5]: x_train = train.drop('quality', axis=1) #training df without class column
x_test = test.drop('quality', axis=1) #testing df without class column

y_train = train['quality'] #training df only class column
y_test = test['quality'] #testing df only class column

y_train = y_train.map({3:0, 4:1, 5:2, 6:3, 7:4, 8:5})
y_test = y_test.map({3:0, 4:1, 5:2, 6:3, 7:4, 8:5})
```

```
In [6]: n_inputs = [x_train.shape[1]] #n cols => n inputs in model
n_units = 12
n_batch = 100
n_epochs = 10
```

```
In [7]: tf.keras.backend.clear_session() #resets parameters, necessary before each new mode
keras.utils.set_random_seed(0) #setting random seed for the entire program

modelRaw = tf.keras.Sequential([tf.keras.layers.Dense(units=n_units, activation='re
                                tf.keras.layers.Dense(units=6, activation='softmax'
```



```
modelRaw.summary()

modelRaw.compile(optimizer='adam', loss='mae', metrics=['accuracy'])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 12)	60
dense_1 (Dense)	(None, 6)	78
=====		
Total params: 138		
Trainable params: 138		
Non-trainable params: 0		

In [8]: `modelRaw.fit(x_train, y_train, batch_size=n_batch, epochs=n_epochs)`

```
Epoch 1/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 2/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 3/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 4/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 5/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 6/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 7/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 8/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 9/10
13/13 [=====] - 0s 1ms/step - loss: 2.4726 - accuracy: 0.39
87
Epoch 10/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.39
87
```

Out[8]: `<keras.callbacks.History at 0x26ca4736f40>`

In [9]: `###end raw`

`###begin scaled`

In [10]: `#recombining train & test to get overall max and min values so test and train are s`
`whole_set = pd.concat([train,test])`
`whole_set.describe() #summary to show all columns have varying scales`

Out[10]:

	volatile acidity	citric acid	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.527821	0.270976	0.658149	10.422983	5.636023
std	0.179060	0.194801	0.169507	1.065668	0.807569
min	0.120000	0.000000	0.330000	8.400000	3.000000
25%	0.390000	0.090000	0.550000	9.500000	5.000000
50%	0.520000	0.260000	0.620000	10.200000	6.000000
75%	0.640000	0.420000	0.730000	11.100000	6.000000
max	1.580000	1.000000	2.000000	14.900000	8.000000

In [11]: `#drop quality class label column before scaling`
`whole_set.drop('quality',axis=1, inplace=True)`

`#build scaler`
`scaler = MinMaxScaler() #build scaler`
`scaler.fit(whole_set) #fit scaler to entire df w/o quality col`

Out[11]: `MinMaxScaler()`

In [12]: `x_trainScaled=scaler.transform(x_train)`
`x_testScaled=scaler.transform(x_test)`

`#make transformed data in a dataframe (.transform returns arrays, we want df) using`
`x_trainScaled = pd.DataFrame(x_trainScaled, columns=x_train.columns)`
`x_testScaled = pd.DataFrame(x_testScaled, columns=x_test.columns)`

In [13]: `x_trainScaled.describe() #now each attr column has min 0 and max 1`

Out[13]:

	volatile acidity	citric acid	sulphates	alcohol
count	1279.000000	1279.000000	1279.000000	1279.000000
mean	0.281668	0.266059	0.195751	0.310976
std	0.120700	0.193606	0.099773	0.166797
min	0.000000	0.000000	0.000000	0.000000
25%	0.191781	0.090000	0.131737	0.169231
50%	0.273973	0.250000	0.173653	0.276923
75%	0.356164	0.420000	0.239521	0.415385
max	1.000000	1.000000	1.000000	1.000000

```

In [14]: tf.keras.backend.clear_session() #resets parameters, necessary before each new mode

modelScaled = tf.keras.Sequential([tf.keras.layers.Dense(units=n_units, activation=
                                tf.keras.layers.Dense(units=6, activation='soft

modelScaled.summary()

modelScaled.compile(optimizer='adam', loss='mae', metrics=['accuracy'])

modelScaled.fit(x_trainScaled, y_train, batch_size=n_batch, epochs=n_epochs)

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	60
dense_1 (Dense)	(None, 6)	78

=====
 Total params: 138
 Trainable params: 138
 Non-trainable params: 0


```
Epoch 1/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 2/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 3/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 4/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 5/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 6/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 7/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 8/10
13/13 [=====] - 0s 1ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 9/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
Epoch 10/10
13/13 [=====] - 0s 2ms/step - loss: 2.4726 - accuracy: 0.10
16
```

Out[14]: <keras.callbacks.History at 0x26ca5aa44f0>

```
In [15]: ###end scaled

        ###begin oversampled
```

```
In [16]: #plot to show that oversampling balanced the data
plt.hist(y_train, bins=11)
plt.title("Histogram of Qualities before Oversampling")
plt.ylabel('Number of observations')
plt.xlabel('Quality')
plt.show()
```



```
In [17]: #dividing train into separate dfs for each class value
qual3 = (train[train["quality"]==3])
qual4 = (train[train["quality"]==4])
qual5 = (train[train["quality"]==5])
qual6 = (train[train["quality"]==6])
qual7 = (train[train["quality"]==7])
qual8 = (train[train["quality"]==8])

#number of samples per class label to determine inbalance
n_qual3 = len(qual3) #8
n_qual4 = len(qual4) #42
n_qual5 = len(qual5) #545
n_qual6 = len(qual6) #510
n_qual7 = len(qual7) #159
n_qual8 = len(qual8) #15

n_max = max(n_qual3, n_qual4, n_qual5, n_qual6, n_qual7, n_qual8) #545
```

```
In [18]: #oversample so that each class has the same number of observations,
#equal to the number of observations of quality 5
qual3Oversampled = qual3.sample(n_max, replace=True)
qual4Oversampled = qual4.sample(n_max, replace=True)
qual6Oversampled = qual6.sample(n_max, replace=True)
qual7Oversampled = qual7.sample(n_max, replace=True)
qual8Oversampled = qual8.sample(n_max, replace=True)
```

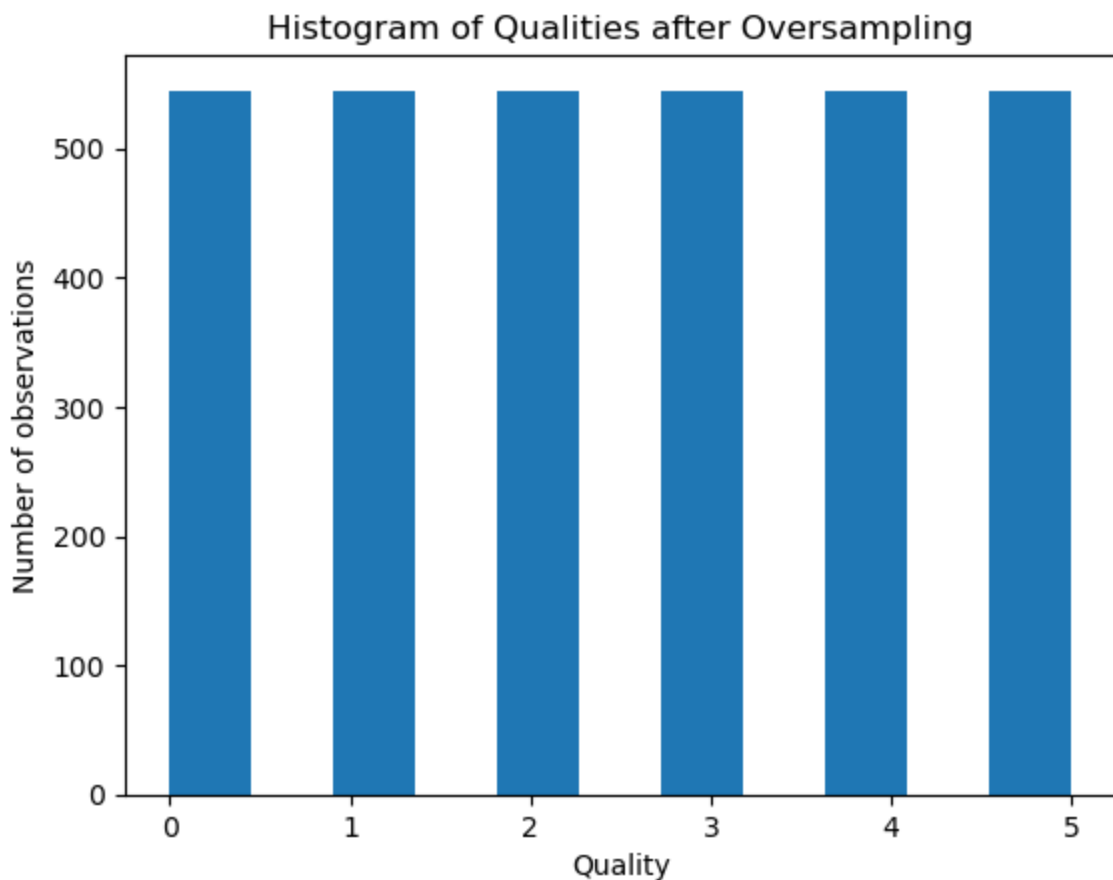
```
In [19]: #concat back into one df. this is unscaled
```

```
trainOversampled = pd.concat([qual30oversampled, qual40oversampled, qual5, qual60oversampled])
```

```
In [20]: #scaling oversampled
x_trainOversampled = trainOversampled.drop('quality',axis=1) #training df without class column
y_trainOversampled = trainOversampled['quality'] #training df only class column

#map values of train from [3,8] to [0,5]
y_trainOversampled = y_trainOversampled.map({3:0, 4:1, 5:2, 6:3, 7:4, 8:5})
```

```
In [21]: #plot to show that oversampling balanced the data
plt.hist(y_trainOversampled, bins=11)
plt.title("Histogram of Qualities after Oversampling")
plt.ylabel('Number of observations')
plt.xlabel('Quality')
plt.show()
```



```
In [22]: #scale trainOversampled w scaler from before
#the same scaler will work bcus our oversampled data will have same range as raw data
x_trainOversampledScaled=scaler.transform(x_trainOversampled)

#make transformed data in a dataframe (.transform returns arrays, we want df) using
x_trainScaledOversampled = pd.DataFrame(x_trainOversampledScaled, columns=x_trainOversampled.columns)
```

```
In [23]: tf.keras.backend.clear_session() #resets parameters, necessary before each new model
modelOversampledScaled = tf.keras.Sequential([tf.keras.layers.Dense(units=n_units,
                                                                    activation='relu',
                                                                    input_shape=(x_trainScaledOversampled.shape[1],)),
                                              tf.keras.layers.Dense(units=6, activation='softmax')])
```

```
modelOversampledScaled.summary()

modelOversampledScaled.compile(optimizer='adam', loss='mae', metrics=['accuracy'])

modelOversampledScaled.fit(x_trainOversampledScaled, y_trainOversampled, batch_size
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 12)	60
dense_1 (Dense)	(None, 6)	78

=====

Total params: 138
Trainable params: 138
Non-trainable params: 0

Epoch 1/10
33/33 [=====] - 1s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 2/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 3/10
33/33 [=====] - 0s 1ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 4/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 5/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 6/10
33/33 [=====] - 0s 1ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 7/10
33/33 [=====] - 0s 1ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 8/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 9/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306
Epoch 10/10
33/33 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.1306

Out[23]: <keras.callbacks.History at 0x26ca5d13640>

In [24]: ###end oversampled

```
###begin weighted
#returning to scaled data, then calculating weights and scaling
```

```
In [25]: #make each column a numpy array for class weights computation
qual3_numpy = qual3['quality'].to_numpy()
qual4_numpy = qual4['quality'].to_numpy()
qual5_numpy = qual5['quality'].to_numpy()
qual6_numpy = qual6['quality'].to_numpy()
qual7_numpy = qual7['quality'].to_numpy()
qual8_numpy = qual8['quality'].to_numpy()

#combine numpy arrays into whole numpy and set variable for class values
whole_numpy = np.concatenate((qual3_numpy, qual4_numpy, qual5_numpy, qual6_numpy, q
unique_classes = np.unique(whole_numpy)

#compute weights with sklearn method
weights = sklearn.utils.class_weight.compute_class_weight(class_weight='balanced',
weightsDict = {i:w for i,w in enumerate(weights)})
```

```
In [26]: tf.keras.backend.clear_session() #resets parameters, necessary before each new mode

modelWeighted = tf.keras.Sequential([tf.keras.layers.Dense(units=n_units, activation=
tf.keras.layers.Dense(units=6, activation='softmax')

modelWeighted.summary()

modelWeighted.compile(optimizer='adam', loss='mae', metrics=['accuracy'])

modelWeighted.fit(x_trainScaled, y_train, batch_size=n_batch, epochs=n_epochs, clas
```


Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	60
dense_1 (Dense)	(None, 6)	78

=====
 Total params: 138
 Trainable params: 138
 Non-trainable params: 0

```
Epoch 1/10
13/13 [=====] - 1s 3ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 2/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 3/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 4/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 5/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 6/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 7/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 8/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 9/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
Epoch 10/10
13/13 [=====] - 0s 2ms/step - loss: 2.3889 - accuracy: 0.42
69
```

Out[26]: <keras.callbacks.History at 0x26ca7f117f0>

```
In [27]: #function to output accuracy, precision, recall, and F1 for each model
#input model name, model object, and appropriate test set x_test (scaled or unscaled)
#output header and metric scores as percents

def outputMetrics(modelname, model, x_test):
    y_pred = model.predict(x_test)

    y_pred_df = (pd.DataFrame(y_pred))

    y_pred_label = y_pred.argmax(axis=1)
```

```

f1 = f1_score(y_test,y_pred_label,average="weighted",zero_division=1)
recall = recall_score(y_test,y_pred_label,average="weighted",zero_division=1)
precision = precision_score(y_test, y_pred_label,average="weighted",zero_division=1)
accuracy = accuracy_score(y_test,y_pred_label)
print(modelname, "Metrics:")
print("Accuracy: {:.2f}%".format(accuracy*100))
print("Precision: {:.2f}%".format(precision*100))
print("Recall: {:.2f}%".format(recall*100))
print("F1: {:.2f}%".format(f1*100))

outputMetrics("Raw Model", modelRaw, x_test)
outputMetrics("Scaled Model", modelScaled, x_testScaled)
outputMetrics("Scaled & Oversampled Model", modelOversampledScaled, x_testScaled)
outputMetrics("Weighted Model", modelWeighted, x_testScaled)

```

10/10 [=====] - 0s 2ms/step

Raw Model Metrics:

Accuracy: 40.00%

Precision: 76.00%

Recall: 40.00%

F1: 22.86%

10/10 [=====] - 0s 1ms/step

Scaled Model Metrics:

Accuracy: 5.62%

Precision: 50.71%

Recall: 5.62%

F1: 5.63%

10/10 [=====] - 0s 1ms/step

Scaled & Oversampled Model Metrics:

Accuracy: 3.12%

Precision: 56.69%

Recall: 3.12%

F1: 1.37%

10/10 [=====] - 0s 2ms/step

Weighted Model Metrics:

Accuracy: 42.19%

Precision: 74.86%

Recall: 42.19%

F1: 25.22%

```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
```

```
# Load the wine dataset from a CSV file
```

```
red_train = pd.read_csv('red_wine_train.csv', index_col=0)
```

```
red_test = pd.read_csv('red_wine_test.csv', index_col=0)
```

```
total = pd.concat([red_train, red_test])
```

```
total = total.drop(["quality"], axis=1)
```

```
scaler = StandardScaler()
```

```
scaler.fit(total)
```

```
# Split the data into training and test sets
```

```
X_train = red_train.drop("quality", axis=1)
```

```
X_test = red_train.drop("quality", axis=1)
```

```
y_train = red_train["quality"]
```

```
y_test = red_train["quality"]
```

```
X_train_scaled = scaler.transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
X_train = pd.DataFrame(X_train_scaled, columns=X_train.columns)
```

```
X_test = pd.DataFrame(X_test_scaled, columns=X_test.columns)
```

```
X = pd.concat([X_train, X_test])
```

```
y = pd.concat([y_test, y_train])
```

```
# Train a KNN classifier with k=2 on the training data
```

```
knn = KNeighborsClassifier(n_neighbors=5)
```

```
knn.fit(X_train, y_train)
```

```
# Make predictions on the test data
```

```
y_pred = knn.predict(X_test)
```

```
# Compute the accuracy of the predictions
```

```
prediction_accuracy = accuracy_score(y_test, y_pred)
```

```
print("Prediction Accuracy: {:.2f}%".format(prediction_accuracy*100))
```

```
# Compute cross-validation scores for different values of k
```

```
k_values = [i for i in range(2, 31)]
```

```
scores = []
```

```
for k in k_values:
```

```
    knn = KNeighborsClassifier(n_neighbors=k)
```

```
    score = cross_val_score(knn, X_train, y_train, cv=5)
```

```
    scores.append(np.mean(score))
```

```
# Plot the cross-validation scores vs. k
```

```
plt.plot(k_values, scores, marker='o')
```

```
plt.xlabel("K values")
```

```
plt.ylabel("Accuracy score")
```

```
# Train a KNN classifier with the best value of k on the training data
```

```
knn = KNeighborsClassifier(n_neighbors=22)
```

```
knn.fit(X_train, y_train)
```

```
# Make predictions on the test data
```

```
y_pred = knn.predict(X_test)
```

```
# Compute scores
```

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted', zero_division=1)
recall = recall_score(y_test, y_pred, average='weighted', zero_division=1)
f1 = f1_score(y_test, y_pred, average='weighted', zero_division=1)
```

```
# Print
print("Accuracy: {:.2f}%".format(accuracy*100))
print("Precision: {:.2f}%".format(precision*100))
print("Recall: {:.2f}%".format(recall*100))
print("F1 score: {:.2f}%".format(f1*100))
```

```
plt.show()
```

```
In [18]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

from collections import Counter

from sklearn.metrics import f1_score, accuracy_score, precision_score, recall_score
from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBClassifier
#from hyperopt import STATUS_OK, Trials, fmin, hp, tpe
```

The history saving thread hit an unexpected error (OperationalError('attempt to write a readonly database')).History will not be written to the database.

```
In [19]: red_train = pd.read_csv("data/red_wine_train.csv", index_col=0)
red_test = pd.read_csv("data/red_wine_test.csv", index_col=0)
```

```
In [20]: red_train.describe()
```


```
Out[20]:
```

	fixed acidity	volatile acidity	citric acid	chlorides	total sulfur dioxide	density	su
count	1279.000000	1279.000000	1279.000000	1279.000000	1279.000000	1279.000000	1279
mean	8.282877	0.531235	0.266059	0.087127	46.464816	0.996715	0
std	1.717760	0.176222	0.193606	0.047654	32.173470	0.001916	0
min	4.900000	0.120000	0.000000	0.012000	6.000000	0.990070	0
25%	7.100000	0.400000	0.090000	0.070000	23.000000	0.995545	0
50%	7.900000	0.520000	0.250000	0.079000	38.000000	0.996700	0
75%	9.100000	0.640000	0.420000	0.090000	62.000000	0.997800	0
max	15.900000	1.580000	1.000000	0.611000	278.000000	1.003690	2

First Attempt: outlier removal, and dropping fixed acidity and chlorldies, and the two binary variables (no binning)

```
In [21]: #dropping binary values that were created just
red_train = red_train.drop(["alcohol_higher", "va_high"], axis=1)
red_test = red_test.drop(["alcohol_higher", "va_high"], axis=1)
```

```
In [22]: plt.figure(figsize=(16, 6))
heatmap = sns.heatmap(red_train.corr(), vmin=-1, vmax=1, annot=True, cmap='BrBG')
```



Preprocessing

For decision trees, there typically no reason to do much when it comes to altering the data. Decision trees do not have much problems with scaled, normalized, and other types continous data is not a big deal (especially with XGB). If we had categorical data, this would be different, espically if we had to data cleanup, but other than outliers there is not much that can more be done that what is given.

```
In [23]: def remove_outliers(df: pd.DataFrame, n :float, columns):
    #this is the Tukey ruel which gets the values that exists outside of the outer
    #This is valuable if outliers effect the data alot (typically regression), in c
    #From EDA it seems that we should not remove outliers as they can help point to
    total_outliers = []

    for col in columns:

        #generating the quantile ranges that will be used to determine outliers
        q1 = df[col].quantile(.25)
        q3 = df[col].quantile(.75)

        iqr = q3 - q1
        outer_fence = iqr * 1.5

        outliers = df[(df[col] < q1 - outer_fence) | (df[col] > q3 + outer_fence)].

        total_outliers.extend(outliers)

    #select the indexes (tuples) that have more than n attributes that are outliers
    #creates an object that has keys (index), with values (amount of apperences, wh
    outliers = Counter(total_outliers)

    #iterates over all items and reutrns the
    items_greater =[]
    for i in outliers.items():
        if(i[1] >= n):
            items_greater.append(i[0])
```

```
return items_greater
```

```
In [24]: #chose two as during Eda it seemed that individually there was not much outliers an
print("old len: " + str(len(red_train)))
outliers = remove_outliers(red_train, 2 , red_train.columns[:-1])
print("new len: " + str(len(red_train.drop(outliers, axis = 0).reset_index(drop=True))))
```

old len: 1279

new len: 1213

overall outliers are negligible

Model Testing

```
In [25]: X_train = red_train.drop("quality" , axis=1)
y_train = red_train["quality"]
X_test = red_test.drop("quality" , axis=1).copy()
y_test = red_test["quality"]
```

```
In [26]: #xgb expects only what is present, not the actual scale of the data relative to qua
#therefore, we have to map to values 0-5 for the scale
y_train = y_train.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5})
y_test = y_test.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5})
```

Hyperopt

-- uses bayesian optimization to find the best parameter for machine learning algorithm, by using probabilistic search of the hyperparameters supplied

Compared to exhaustive search it is much faster and its performance is only a bit lower

How to implement:

1. initialise the domain space (same as a grid search)
2. define the objective function that we want to minimize (error rate) of the model that we are testing (XGBoost decision trees in this case)
3. Optimize algorithm choice (the method used to construct the surrogate objective function)
4. Results, the score or the value pairs that the algorithm uses to build the model

Below the hyperopt stuff that I personally was messing around with, but I do not think that I am going to include it in final report. Just skip...

```
In [27]: # space = {"max_depth":hp.quniform("max_depth",3,18,1), #the max dept with integer
#           "gamma" : hp.uniform('gamma',1,9), #the gamma values with values ranging
#           'colsample_bytree' : hp.uniform('colsample_bytree', 0.5,1), #ratio of col
#           'min_child_weight' : hp.quniform('min_child_weight', 0, 10, 1), #tuning t
#           'n_estimators': 180,
```

```

#         'seed': 0
#     }

# def ojective(space:dict):
#     #creating a classifier with the opametrs pulled from the space that has be
#     #Most explanation is above
#     clf=XGBClassifier(
#         n_estimators =space['n_estimators'],
#         max_depth = int(space['max_depth']),
#         gamma = space['gamma'],
#         min_child_weight=int(space['min_child_weight']),
#         colsample_bytree=int(space['colsample_bytree']),
#         objective="multi:softprob", #type of objective function that is u
#         early_stopping_rounds=10, #sets the early stopping rounds if the
#         eval_metric="auc") #the measure to determine within the gradient

#     #evaluation train set and test set for doing a fit, effcivly measuring
#     evaluation = [( X_train, y_train), ( X_test, y_test)]

#     clf.fit(X_train,
#             y_train,
#             eval_set=evaluation, #passed in for the set
#             verbose=False)

#     # make a prediction
#     y_pred = clf.predict(X_test)

#     accuracy = f1_score(y_test,y_pred,average="weighted",zero_division=1)
#     print("SCORE:" + str(accuracy))
#     return {'loss' : -accuracy, "status" : STATUS_OK}

# trials = Trials()

# best_hyperparams = fmin(fn = ojective,
#                          space = space,
#                          algo = tpe.suggest,
#                          max_evals = 50,
#                          trials = trials)

```

Available optimizaiton algorithms: (oudated)

- hp.choice(label,options) : returns a choice of one of the options
- hp.randint(label,upper) : returns a random integer better range of 0 --> upper
- hp.uniform(label,low, high) : returns a value uniformly between the low and high
- hp.uniform(label,low,high,q) : returns a value round to (uniform(low,high)/q) , and returns an integer


```
-- hp.normal(label,mean,std) : returns a real value that is
normally distributed with mean and standard deviation
```

Trials:

```
-- an object that contains or stores all the relevent information such as a hyperparameter.
The loss functions for each type of parameter is stored here. Whenever doing iterations of
training with current hyperparameters.
```

```
-- fmin is an optimization function that minimizes the loss function for each paramter inside
of space.
```

```
-- algo: The type of the algorithm for finding best hyperparamter. Tpe is a type of decision
tree, so effectively we are using a decision tree to do the hyperparamter choosing.
```

```
-- max evals: the amount of iterations that we choose to run through
```

Hyper parameters:

```
Booster: choose the type of booster to use (we will use tree in
this case)
```

```
-- 3 options
tree( gbtree,dart)
linear(gblinear)
```

```
Booster Parameters: only the tree booster ones, only listing the
ones usuful to mulitlable imbalanced data (trees for the win)
```

```
-- eta: the learning rat for Gradient boosting, and its range
typically is 0.01 - 0.2
-- gamma: how the node is split in a tree, the larger the more
conservative a tree is, range(0 --> infinity)
-- max_depth: maximum depth of a tree typical values are (3-
10), should use cv
-- min_child_weight: tune using cv but range is 0-->infinite
-- subsample: fraction of observations to be samples for tree,
lower values more conservative, typical values (0,1)
-- colsample_bytree: ratio of columns when construction each
tree
-- colsample_bylevel: ratio of columns at each level of the
tree
-- tree method: constuction algorithm used in model (multiple
choices)
-- max_leaves: is maximum number of nodes to be added
```

Others

```
-- alpha : used for lasso regression, increasing makes the
model more conservative
-- lambda : used for ridge regression, increasing makes the
model more conservative
```

Learning Task: parameters used to define the optimization objective for learning

--objective: should use multi:softprob or multi: softmax
 --eval metric: should use auc, or merror

```
In [28]: space = {'max_depth': [3, 6, 10, 15, 20],
                  'learning_rate': [0.01, 0.1, 0.2, 0.3, 0.4],
                  'subsample': np.arange(0.5, 1.0, 0.1),
                  'gamma' : np.arange(1,9,1),
                  'colsample_bytree': np.arange(0.5, 1.0, 0.1),
                  'colsample_bylevel': np.arange(0.5, 1.0, 0.1),
                  'min_child_weight' : np.arange(1, 10, 1),
                  'n_estimators': [100,180, 250, 500, 750],
                  }
```

```
In [29]: #making the base model
#multi prob is a vector, containing all the classes
model = XGBClassifier(objective="multi:softprob",eval_metric="auc",)
clf = RandomizedSearchCV(estimator=model,
                        param_distributions=space, #assigning space
                        scoring="f1_weighted", #eval metric for the hyperparms, we
                        n_iter=25, #amount of iterations per cv (random combinatio
                        n_jobs=4, #amount of parralel processes to run
                        random_state=1)

clf.fit(X_train,y_train)
```

```
Out[29]: RandomizedSearchCV
          estimator: XGBClassifier
              XGBClassifier
```

```
In [30]: best_hyperparams = clf.best_params_
best_hyperparams
```

```
Out[30]: {'subsample': 0.8999999999999999,
          'n_estimators': 750,
          'min_child_weight': 9,
          'max_depth': 15,
          'learning_rate': 0.1,
          'gamma': 2,
          'colsample_bytree': 0.8999999999999999,
          'colsample_bylevel': 0.6}
```

```
In [31]: accuracy_f1 = []
recall = []
precision = []
accuracy = []
for i in range(0,5):
    #creating model with best hyperparameters
    clf=XGBClassifier(
        n_estimators = best_hyperparams['n_estimators'],
```

```

max_depth = int(best_hyperparams['max_depth']),
learning_rate = best_hyperparams['learning_rate'],
gamma = best_hyperparams['gamma'],
min_child_weight=int(best_hyperparams['min_child_weight']),
colsample_bytree=int(best_hyperparams['colsample_bytree']),
colsample_bylevel=best_hyperparams['colsample_bylevel'],
objective="multi:softprob", #type of objective function that is used, you h
eval_metric="auc", #the measure to determine within the gradient boosting t
seed=i)

clf.fit(X_train,y_train)

y_pred = clf.predict(X_test)

#appending the scores from the particular rune to the list
accuracy_f1.append(f1_score(y_test,y_pred,average="weighted",zero_division=1))
recall.append(recall_score(y_test,y_pred,average="weighted",zero_division=1))
precision.append(precision_score(y_test, y_pred,average="weighted",zero_divisio
accuracy.append(accuracy_score(y_test,y_pred))

```

```

In [32]: average_f1 = np.mean(accuracy_f1)
average_recall = np.mean(recall)
average_precision = np.mean(precision)
average_accuracy = np.mean(accuracy)
print("Average F1: " + str(average_f1))
print("Average Recall: " + str(average_recall))
print("Average Precision: " + str(average_precision))
print("Average Accuracy: " + str(average_accuracy))

```

```

Average F1: 0.5943817969672713
Average Recall: 0.62375
Average Precision: 0.653822962237242
Average Accuracy: 0.62375

```

This file also contains the EDA and Preprocessing, but since this came after doing the other attempts, and it has overall lower correlation with data, we decided that only one locations should suffice.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

from collections import Counter

from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.metrics import f1_score, accuracy_score, recall_score, precision_score
from xgboost import XGBClassifier
#from hyperopt import STATUS_OK, Trials, fmin, hp, tpe
```

The history saving thread hit an unexpected error (OperationalError('attempt to write a readonly database')).History will not be written to the database.


```
In [2]: #this was used to randomize the data as it was made by just adding one of wine file
"""
total_wine = pd.read_csv("data/winequality-total.csv", delimiter=";")
total = total_wine.sample(frac=1)
total_wine.to_csv("data/winequality-total.csv")
"""
```

```
Out[2]: '\ntotal_wine = pd.read_csv("data/winequality-total.csv", delimiter=";")\ntotal =
total_wine.sample(frac=1)\ntotal_wine.to_csv("data/winequality-total.csv")\n'
```

```
In [3]: total_wine = pd.read_csv("data/winequality-total.csv", index_col=0)
```

```
In [4]: plt.figure(figsize=(16,6))
sns.heatmap(total_wine.corr()[["quality"]].sort_values(by='quality', ascending=False
```

```
Out[4]: <Axes: >
```



The only valuable ones that seem to be present: alcohol, volatile acidity, chlorides, and density. However, to maintain consistence, and to also test the effectiveness of correlation analysis, we will drop the same ones we did for the xgboost implementation on only red wine. Granted, the dynamics do change with the entrance of white wine, but if the accuracy is at least comparable, we can investigate later.

```
In [5]: total_wine = total_wine.drop(["residual sugar", "pH", "free sulfur dioxide"], axis=1)
```

Alcohol

```
In [6]: fig = make_subplots(rows=2, cols=2)

#histogram for red
fig.append_trace(go.Histogram(
    x=total_wine["alcohol"],
    name="red hist"), row=1, col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=total_wine["alcohol"],
    name="red box"
), row=1, col=2)

fig.append_trace(go.Box(
    x=total_wine["quality"],
    y=total_wine["alcohol"],
    name="quality/alcohol"
), row=2, col=1)

#histogram for white
fig.update_layout(height=600, width=1200, title_text="Alcohol data")
fig.show()
```

Allmost the same as what was seen ealier in red_wine. Therefore, going to treat it about the same


Volatile acidity

```
In [7]: fig = make_subplots(rows=2,cols=2)

#histogram for red
fig.append_trace(go.Histogram(
    x=total_wine["volatile acidity"],
    name="normal hist"), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=total_wine["volatile acidity"],
    name="normal box"
),row=1,col=2)
```

```
fig.append_trace(go.Box(  
    x=total_wine["quality"],  
    y=total_wine["volatile acidity"],  
    name="quality/alcohol"  
),row=2, col=1)  
  
#histogram for white  
fig.update_layout(height=600, width=1200, title_text="VA data")  
fig.show()
```



Extremely right skewed, can be adjusted later if necessary, but it seems that the lower the value goes the more likely it is high quality, but not a strong indicator

Density

```
In [8]: fig = make_subplots(rows=2,cols=2)

#histogram for red
fig.append_trace(go.Histogram(
    x=total_wine["density"],
    name="normal hist"), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=total_wine["density"],
    name="normal box"
),row=1,col=2)

fig.append_trace(go.Box(
    x=total_wine["quality"],
    y=total_wine["density"],
    name="quality/density"
),row=2, col=1)

#histogram for white
fig.update_layout(height=600, width=1200, title_text="density data")
fig.show()
```


Should drop that outlier but follow a similar trend to VA, where the values tend to decrease. Additionally, it appears to still be a shaky, but still gaussian distribution.

Chlorides

```
In [9]: fig = make_subplots(rows=2,cols=2)

#histogram for red
fig.append_trace(go.Histogram(
    x=total_wine["chlorides"],
    name="normal hist"), row=1,col=1)

#boxplot for red
fig.append_trace(go.Box(
    x=total_wine["chlorides"],
```

```
        name="normal box"
    ),row=1,col=2)

fig.append_trace(go.Box(
    x=total_wine["quality"],
    y=total_wine["chlorides"],
    name="quality/chlorides"
),row=2, col=1)

#histogram for white
fig.update_layout(height=600, width=1200, title_text="chlorides data")
fig.show()
```

Very far skew, but it seems relevant. The

Preprocessing

```
In [10]: def remove_outliers(df: pd.DataFrame, n :float, columns):
    #this is the Tukey ruel which gets the values that exists outside of the outer
    #This is valuable if outliers effect the data alot (typically regression), in c
    #From EDA it seems that we should not remove outliers as they can help point to
    total_outliers = []

    for col in columns:

        #generating the quantile ranges that will be used to determine outliers
        q1 = df[col].quantile(.25)
        q3 = df[col].quantile(.75)

        iqr = q3 - q1
        outer_fence = iqr * 1.5

        outliers = df[(df[col] < q1 - outer_fence) | (df[col] > q3 + outer_fence)].

        total_outliers.extend(outliers)

    #select the indexes (tuples) that have more than n attributes that are outliers
    #creates an object that has keys (index), with values (amount of apperences, wh
    outliers = Counter(total_outliers)

    #iterates over all items and reutrns the
    items_greater = []
    for i in outliers.items():
        if(i[1] >= n):
            items_greater.append(i[0])

    return items_greater
```

```
In [11]: len(total_wine)
```

```
Out[11]: 6497
```

```
In [12]: #finding all outliers that have significant outliers
outliers = remove_outliers(total_wine,2,total_wine.columns[:-1])
len(total_wine.drop(outliers, axis=0).reset_index(drop=True))
```

```
Out[12]: 6109
```

Since it is so low and Dt are natually resistant, am going to keep outliers in (might hugley effect our imbalanced)

Training Model (with same features as red wine)

```
In [13]: total_train, total_test = train_test_split(total_wine, test_size=0.2, stratify=tota
```

```
In [14]: X_train = total_train.drop("quality" , axis=1)
y_train = total_train["quality"]
```

```
X_test = total_test.drop("quality" , axis=1).copy()
y_test = total_test["quality"]
```

```
In [15]: #need to map the data for XGBoost to work, expects data in this format
y_train = y_train.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5, 9:6})
y_test = y_test.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5, 9:6})
```

Hyper parameters:

Booster: choose the type of booster to use (we will use tree in this case)

```
-- 3 options
tree( gbtree,dart)
linear(gblinear)
```

Booster Parameters: only the tree booster ones, only listing the ones useful to multitable imbalanced data (trees for the win)

```
-- eta: the learning rat for Gradient boosting, and its range typically is 0.01 - 0.2
-- gamma: how the node is split in a tree, the larger the more conservative a tree is, range(0 --> infinity)
-- max_depth: maximum depth of a tree typical values are (3-10), should use cv
-- min_child_weight: tune using cv but range is 0-->infinite
-- subsample: fraction of observations to be samples for tree, lower values more conservative, typical values (0,1)
-- colsample_bytree: ratio of columns when construction each tree
-- colsample_bylevel: ratio of columns at each level of the tree
-- tree method: constuction algorithm used in model (multiple choices)
-- max_leaves: is maximum number of nodes to be added
```

Others

```
-- alpha : used for lasso regression, increasing makes the model more conservative
-- lambda : used for ridge regression, increasing makes the model more conservative
```

Learning Task: parameters used to define the optimization objective for learning

```
--objective: should use multi:softprob or multi: softmax
--eval metric: should use auc, or merror
```

```
In [16]: space = {'max_depth': [3, 6, 10, 15, 20],
                  'learning_rate': [0.01, 0.1, 0.2, 0.3, 0.4],
                  'subsample': np.arange(0.5, 1.0, 0.1),
                  'gamma' : np.arange(1,9,0.1),
                  'colsample_bytree': np.arange(0.5, 1.0, 0.1),
```

```
'colsample_bylevel': np.arange(0.5, 1.0, 0.1),
'min_child_weight' : np.arange(1, 10, 1),
'n_estimators': [100,150, 250, 500, 750],
}
```

```
In [17]: #making the base model
#multi_prob is a vector, containing all the classes a
model = XGBClassifier(objective="multi:softprob",eval_metric="auc",)
clf = RandomizedSearchCV(estimator=model,
                        param_distributions=space, #assigning space
                        scoring="f1_weighted", #eval metric for the hyperparams, we
                        n_iter=25, #amount of iterations per cv (random combination)
                        n_jobs=4, #amount of parallel processes to run
                        random_state=1)

clf.fit(X_train,y_train)
```

/home/cole/anaconda3/envs/datasci/lib/python3.10/site-packages/sklearn/model_selection/_split.py:700: UserWarning:

The least populated class in y has only 4 members, which is less than n_splits=5.

```
Out[17]: RandomizedSearchCV
          estimator: XGBClassifier
              XGBClassifier
```

```
In [18]: #assigning best params
best_hyperparams = clf.best_params_
best_hyperparams
```

```
Out[18]: {'subsample': 0.7999999999999999,
          'n_estimators': 250,
          'min_child_weight': 4,
          'max_depth': 15,
          'learning_rate': 0.1,
          'gamma': 1.7000000000000006,
          'colsample_bytree': 0.6,
          'colsample_bylevel': 0.5}
```

```
In [19]: accuracy_f1 = []
recall = []
precision = []
accuracy = []
for i in range(0,5):
    clf=XGBClassifier(
        n_estimators = best_hyperparams['n_estimators'],
        max_depth = int(best_hyperparams['max_depth']),
        learning_rate = best_hyperparams['learning_rate'],
        gamma = best_hyperparams['gamma'],
        min_child_weight=int(best_hyperparams['min_child_weight']),
        colsample_bytree=int(best_hyperparams['colsample_bytree']),
        colsample_bylevel=best_hyperparams['colsample_bylevel'],
```

```

objective="multi:softprob", #type of objective function that is used, you h
eval_metric="auc", #the measure to determine within the gradient boosting t
seed=i)

clf.fit(X_train,y_train)

y_pred = clf.predict(X_test)

#appending the scores from the particular rune to the list
accuracy_f1.append(f1_score(y_test,y_pred,average="weighted",zero_division=1))
recall.append(recall_score(y_test,y_pred,average="weighted",zero_division=1))
precision.append(precision_score(y_test, y_pred,average="weighted",zero_divisio
accuracy.append(accuracy_score(y_test,y_pred))

```

```

In [20]: average_f1 = np.mean(accuracy_f1)
average_recall = np.mean(recall)
average_precision = np.mean(precision)
average_accuracy = np.mean(accuracy)
print("Average F1: " + str(average_f1))
print("Average Recall: " + str(average_recall))
print("Average Precision: " + str(average_precision))
print("Average Accuracy: " + str(average_accuracy))

```

Average F1: 0.5273647702420745

Average Recall: 0.558

Average Precision: 0.5370645680886247

Average Accuracy: 0.558


Testing Model (only using high direct corelation)

```

In [21]: plt.figure(figsize=(16,6))
sns.heatmap(total_wine.corr()[["quality"]].sort_values(by='quality',ascending=False)

```

Out[21]: <Axes: >



```

In [22]: #dropping all values that are below 0.1 and between -0.1
total_wine_dropped = total_wine.drop(["citric acid", "sulphates", "total sulfur dio

```

```
In [23]: X_train = total_wine_dropped.drop("quality" , axis=1)
y_train = total_wine_dropped["quality"]
X_test = total_wine_dropped.drop("quality" , axis=1).copy()
y_test = total_wine_dropped["quality"]
```

```
In [24]: #mapping values to different type for algorithm
y_train = y_train.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5, 9:6})
y_test = y_test.map({3: 0, 4: 1, 5:2, 6:3, 7:4, 8:5, 9:6})
```

```
In [26]: #making the base model
model = XGBClassifier(objective="multi:softprob",eval_metric="auc",)
clf = RandomizedSearchCV(estimator=model,
                        param_distributions=space,
                        scoring="f1_weighted",
                        n_iter=25,
                        n_jobs=4,
                        random_state=1)

clf.fit(X_train,y_train)
```

```
Out[26]: RandomizedSearchCV
  estimator: XGBClassifier
    XGBClassifier
```

```
In [27]: best_hyperparams = clf.best_params_
best_hyperparams
```

```
Out[27]: {'subsample': 0.6,
'n_estimators': 150,
'min_child_weight': 1,
'max_depth': 15,
'learning_rate': 0.2,
'gamma': 6.2000000000000005,
'colsample_bytree': 0.6,
'colsample_bylevel': 0.5}
```

```
In [28]: accuracy_f1 = []
recall = []
precision = []
accuracy = []
for i in range(0,5):
    clf=XGBClassifier(
        n_estimators = best_hyperparams['n_estimators'],
        max_depth = int(best_hyperparams['max_depth']),
        learning_rate = best_hyperparams['learning_rate'],
        gamma = best_hyperparams['gamma'],
        min_child_weight=int(best_hyperparams['min_child_weight']),
        colsample_bytree=int(best_hyperparams['colsample_bytree']),
        colsample_bylevel=best_hyperparams['colsample_bylevel'],
        objective="multi:softprob", #type of objective function that is used, you h
        eval_metric="auc", #the measure to determine within the gradient boosting t
        seed=i)
```

```
clf.fit(X_train,y_train)

y_pred = clf.predict(X_test)

#appending the scores from the particular rune to the list
accuracy_f1.append(f1_score(y_test,y_pred,average="weighted",zero_division=1))
recall.append(recall_score(y_test,y_pred,average="weighted",zero_division=1))
precision.append(precision_score(y_test, y_pred,average="weighted",zero_division=1))
accuracy.append(accuracy_score(y_test,y_pred))
```

```
In [29]: average_f1 = np.mean(accuracy_f1)
average_recall = np.mean(recall)
average_precision = np.mean(precision)
average_accuracy = np.mean(accuracy)
print("Average F1: " + str(average_f1))
print("Average Recall: " + str(average_recall))
print("Average Precision: " + str(average_precision))
print("Average Accuracy: " + str(average_accuracy))
```

```
Average F1: 0.5142225228156577
Average Recall: 0.5545636447591196
Average Precision: 0.5883452301824906
Average Accuracy: 0.5545636447591196
```